



1 A fundamental trade-off in VLM inference

A high-resolution visual question




Q: What time does the clock show?
A. Three o'clock. B. Two o'clock. ✓
C. Nine o'clock. D. Twelve o'clock.

Option A: More visual tokens, smaller backbone




Same inference budget
Which allocation works better?

Option B: Fewer visual tokens, larger backbone




2 Does scaling resolve every error?

Questions that benefit from scaling



Performance
**improves with
scale** for many
items.

Questions that remain unresolved



A substantial
portion **remains
largely incorrect**
over the tested
range.