




Higher resolution

Finer visual detail

Visual tokens R



Larger backbone

Greater reasoning capacity

Language backbone size N

Given a fixed inference budget
How should compute be split
between N and R ?

Where does scaling stop?
Which questions remain hard?

How does performance scale?
Do architectures differ?

How should compute be split?
What should we deploy?

$$L(N, R) = e_{\infty} + \frac{A}{N^{\alpha}} + \frac{B}{R^{\beta}}$$

Separable law

InternVL and *QwenVL*
show different scaling behavior

Compute
allocation

Cost model + inference budget

Deployment
choice